

# Linear Regression & Logistic Regression

# Machine Learning



- Machine learning means building algorithms based on a **set of examples** from the same phenomenon.
- Machine learning can also be seen as a process for solving real-world problems in considered that it consists of two parts:
  1. Data collection
  2. Building a statistical model based on data

# Types of learning



- Supervised learning
- Unsupervised learning
- Reinforcement learning


# Supervised Learning



- The dataset is the collection of **labeled** examples;  $\{(x_i, y_i)\}_{i=1, \dots, N}$
- **Goal:** Learn a function to map  $x \rightarrow y$
- Each element  $x_i$  among  $N$  is called a feature vector.
- A **feature** vector is a vector in which each dimension contains a value that describes the example somehow.
- The label  $y_i$  can be either an element belonging to
  - a finite set of classes  $\{1, 2, \dots, C\}$ ,
  - a real number,
  - a more complex structure, like a vector, a matrix, a tree, or a graph.


# Unsupervised Learning

- In unsupervised learning, the dataset is a collection of **unlabeled** examples,  $\{x_i\}_{i=1,\dots,N}$  (No label!)
- **Goal:** Learn some underlying hidden structure of the data that can be used to solve a practical problem.



# Reinforcement Learning

- Problems involving an **agent** interacting with an **environment**, which provides numeric **reward** signals
- **Goal:** Learn how to take actions in order to maximize reward



# Types of learning




- Supervised learning. Examples:
  - Regression
    - In regression the output is **continuous**
  - Classification
    - For classification the output(s) is nominal
  - object detection
  - image captioning
  - etc
- Unsupervised learning
- Reinforcement learning

# Regression

- In regression the output is **continuous**
- Function Approximation
  - Many models could be used – Simplest is linear regression
  - We are looking to build a linear model as follows:

$$f_{w,b}(\mathbf{x}) = \mathbf{w}\mathbf{x} + \mathbf{b}$$

$y$   
dependent  
variable  
(output)




$x$  – independent variable (input)



# Linear Regression

- To find the best  $\mathbf{w}$  and  $\mathbf{b}$ , the optimization problem that needs to be solved come as follows:
- Minimize the cost function

$$\frac{1}{N} \sum_{i=1}^N \underbrace{(f_{w,b}(\mathbf{x}_i) - y_i)^2}_{\text{Loss function}}$$



# Two approaches:



1. Direct solution
2. Gradient descent

Cost function:

$$\begin{aligned} J(w, b) &= \frac{1}{N} \sum_{i=1}^N L(y^{(i)}, t^{(i)}) \\ &= \frac{1}{2N} \sum_{i=1}^N (y^{(i)} - t^{(i)})^2 \end{aligned}$$

# Two approaches:



1. Direct solution
2. Gradient descent

Cost function:

$$\begin{aligned} J(w, b) &= \frac{1}{N} \sum_{i=1}^N L(y^{(i)}, t^{(i)}) \\ &= \frac{1}{2N} \sum_{i=1}^N (y^{(i)} - t^{(i)})^2 \end{aligned}$$

Derivatives:

$$\begin{aligned} \frac{\partial J}{\partial w_j} &= \frac{1}{N} \sum_{i=1}^N (y^{(i)} - t^{(i)}) x_j^{(i)} \\ \frac{\partial J}{\partial b} &= \frac{1}{N} \sum_{i=1}^N (y^{(i)} - t^{(i)}) \end{aligned}$$

# Direct solution

- Chain rule for derivatives:

$$\begin{aligned}\frac{\partial \mathcal{L}}{\partial w_j} &= \frac{d\mathcal{L}}{dy} \frac{\partial y}{\partial w_j} \\ &= \frac{d}{dy} \left[ \frac{1}{2}(y - t)^2 \right] \cdot x_j \\ &= (y - t)x_j \\ \frac{\partial \mathcal{L}}{\partial b} &= y - t\end{aligned}$$

- Cost derivatives (average over data points):

$$\begin{aligned}\frac{\partial \mathcal{E}}{\partial w_j} &= \frac{1}{N} \sum_{i=1}^N (y^{(i)} - t^{(i)}) x_j^{(i)} \\ \frac{\partial \mathcal{E}}{\partial b} &= \frac{1}{N} \sum_{i=1}^N y^{(i)} - t^{(i)}\end{aligned}$$

# Direct solution



- The minimum must occur at a point where the partial derivatives are zero.


$$\frac{\partial \mathcal{E}}{\partial w_j} = 0 \quad \frac{\partial \mathcal{E}}{\partial b} = 0.$$

- This turns out to give a system of linear equations, which we can solve efficiently.
- Optimal weights:

$$\mathbf{w} = (\mathbf{X}^\top \mathbf{X})^{-1} \mathbf{X}^\top \mathbf{t}$$

# Gradient descent

- Gradient descent is an **iterative algorithm**, which means we apply an update repeatedly until some criterion is met.
- We **initialize** the weights to something reasonable (e.g. all zeros) and repeatedly adjust them in the **direction of steepest descent**.



# Gradient descent



- The following update decreases the cost function:

$$\begin{aligned}w_j &\leftarrow w_j - \alpha \frac{\partial \mathcal{E}}{\partial w_j} \\&= w_j - \frac{\alpha}{N} \sum_{i=1}^N (y^{(i)} - t^{(i)}) x_j^{(i)}\end{aligned}$$

- $\alpha$  is a **learning rate**. The larger it is, the faster  $\mathbf{w}$  changes.
  - We'll see later how to tune the learning rate, but values are typically small, e.g. 0.01 or 0.0001

# Multiple features - vectorization

---

- This gets its name from the **gradient**:

$$\frac{\partial \mathcal{E}}{\partial \mathbf{w}} = \begin{pmatrix} \frac{\partial \mathcal{E}}{\partial w_1} \\ \vdots \\ \frac{\partial \mathcal{E}}{\partial w_D} \end{pmatrix}$$

- This is the direction of fastest increase in  $\mathcal{E}$ .
- Update rule in vector form:

$$\begin{aligned} \mathbf{w} &\leftarrow \mathbf{w} - \alpha \frac{\partial \mathcal{E}}{\partial \mathbf{w}} \\ &= \mathbf{w} - \frac{\alpha}{N} \sum_{i=1}^N (y^{(i)} - t^{(i)}) \mathbf{x}^{(i)} \end{aligned}$$



# Gradient descent




- Why gradient descent, if we can find the optimum directly?
  - GD can be applied to a much broader set of models
  - GD can be easier to implement than direct solutions, especially with automatic differentiation software
  - For regression in high-dimensional spaces, GD is more efficient than direct solution (matrix inversion is an  $\mathcal{O}(D^3)$  algorithm).

**Let's implement it ...**

---

# Future studies: Polynomial Regression



---

# Logistic Regression

# Binary Classification



Logistic regression is not really a regression. Rather, it is a learning algorithm for classification.


$y_i$  is a discrete number and indicates the class number. The goal is to build a **linear** model from the labeled data set .

In binary classification is obvious that there is no reason to predict values outside  $[0,1]$ . Let's squash  $y$  into this interval.

# Logistic Activation Function

- The **logistic function** is a kind of **sigmoidal**, or S-shaped, function:

$$\sigma(z) = \frac{1}{1 + e^{-z}}$$



- A linear model with a logistic nonlinearity is known as **log-linear**:

$$z = \mathbf{w}^\top \mathbf{x} + b$$


$$y = \sigma(z)$$

- Used in this way,  $\sigma$  is called an **activation function**, and  $z$  is called the **logit**.

# Cross-entropy Loss

- Because  $y \in [0, 1]$ , we can interpret it as the estimated probability that  $t = 1$ .
- Cross-entropy loss captures this intuition:

$$\mathcal{L}_{\text{CE}}(y, t) = \begin{cases} -\log y & \text{if } t = 1 \\ -\log(1 - y) & \text{if } t = 0 \end{cases}$$
$$= -t \log y - (1 - t) \log(1 - y)$$



# Logistic Regression


— —

$$z = \mathbf{w}^\top \mathbf{x} + b$$

$$y = \sigma(z)$$

$$= \frac{1}{1 + e^{-z}}$$

$$\mathcal{L}_{\text{CE}} = -t \log y - (1 - t) \log(1 - y)$$





# Gradient descent approach



Now to compute the loss derivatives:

$$\begin{aligned}\frac{d\mathcal{L}_{\text{LCE}}}{dz} &= \frac{d}{dz} [t \log(1 + e^{-z}) + (1 - t) \log(1 + e^z)] \\ &= -t \cdot \frac{e^{-z}}{1 + e^{-z}} + (1 - t) \frac{e^z}{1 + e^z} \\ &= -t(1 - y) + (1 - t)y \\ &= y - t\end{aligned}$$

# Comparison of gradient descent updates



- Linear regression:

$$\mathbf{w} \leftarrow \mathbf{w} - \frac{\alpha}{N} \sum_{i=1}^N (y^{(i)} - t^{(i)}) \mathbf{x}^{(i)}$$

- Logistic regression:

$$\mathbf{w} \leftarrow \mathbf{w} - \frac{\alpha}{N} \sum_{i=1}^N (y^{(i)} - t^{(i)}) \mathbf{x}^{(i)}$$